

AWS Sagemaker

In a typical production environment, it is time consuming to deploy machine learning (ML) systems. And it is difficult to deploy and manage environments for ML systems. AWS Sagemaker makes it easier to deploy and manage your ML system environment.

AWS Sagemaker is an end-to-end machine learning platform, which allows to easily deploy production ready machine learning systems. It features auto scaling on deployed distributed ML models. It provides optimized algorithms and also allows for custom built algorithms for use. It provides flexibility in training models by providing support to other frameworks. It supports deep learning frameworks like TensorFlow, Apache MXNet and Gluon. It also supports Apache spark to deploy ML models similar to Spark's MLLib. And custom built algorithms or models based on other frameworks like scikit-learn can be deployed using docker images. It also integrates jupyter notebook for easier writing up of ML model code to be hosted and deployed.

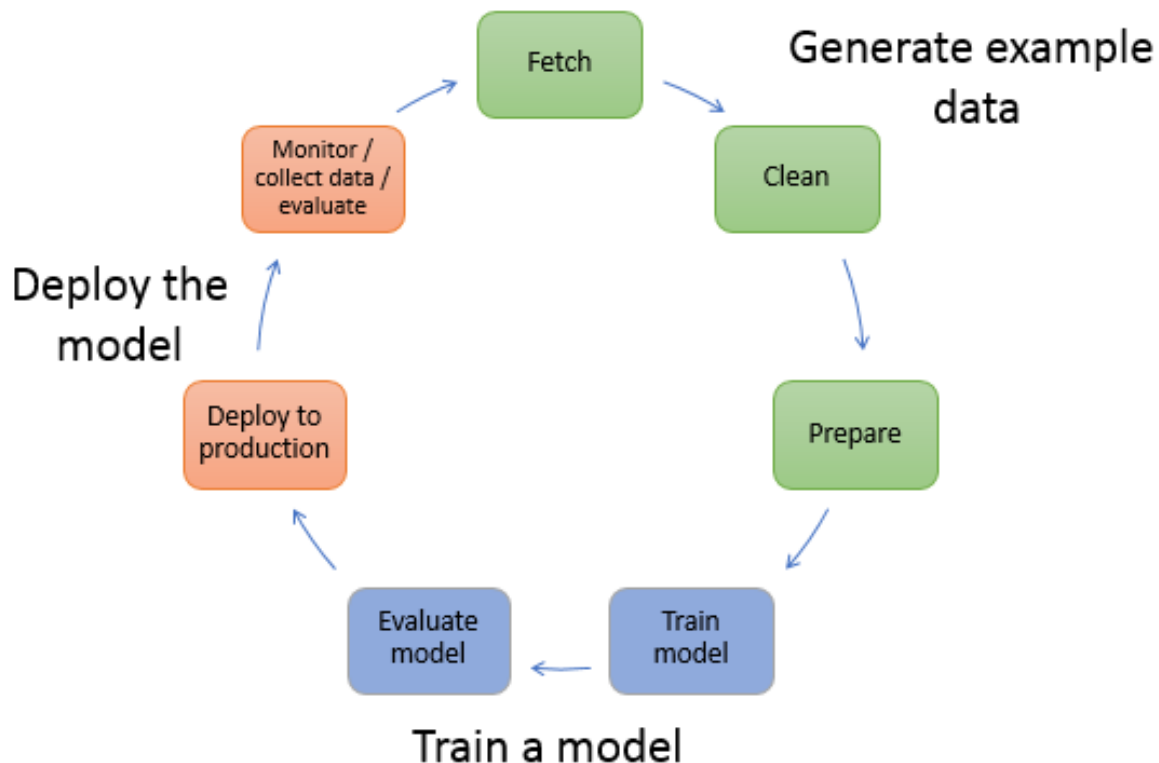
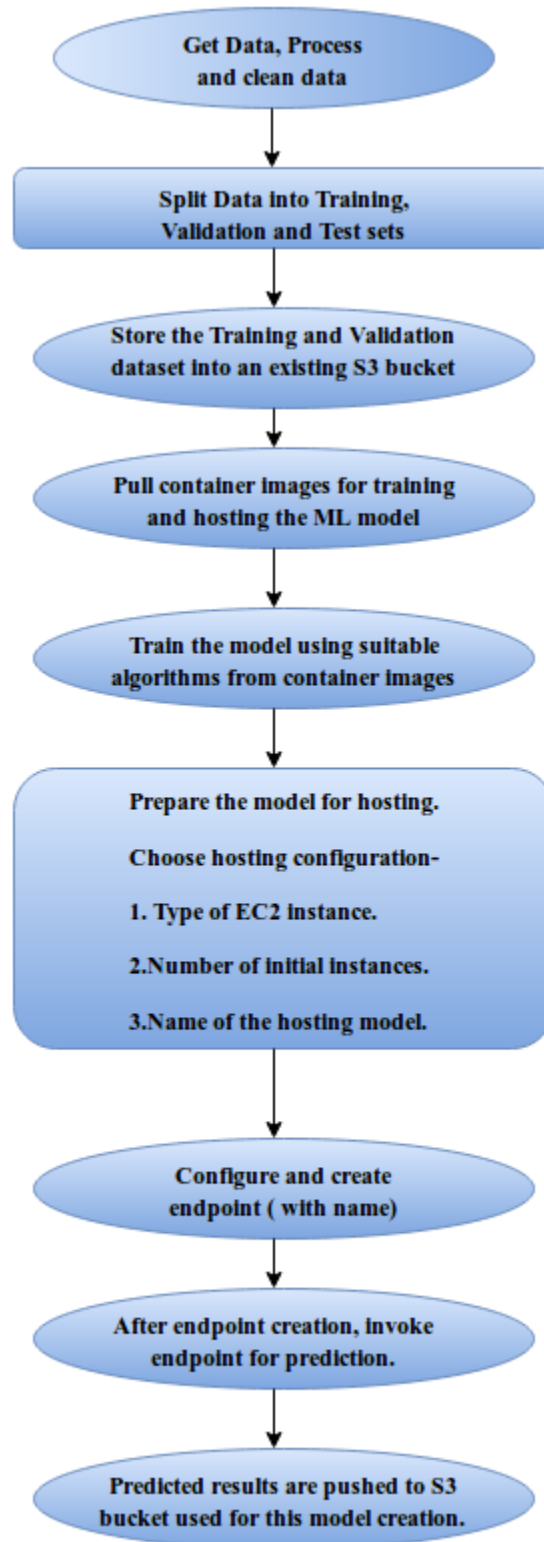


Fig 1. Machine Learning model deployment.

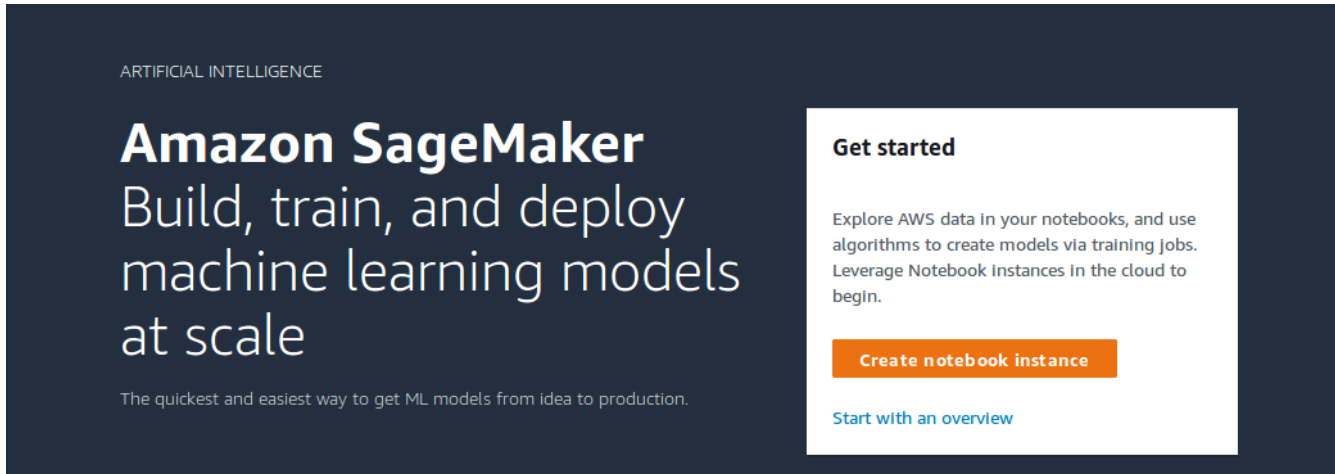
Below flow diagram explains the workflow of AWS Sagemaker.



AWS Sagemaker

Below screenshots are workflow of a sample Model creation on AWS Sagemaker.

1. Create notebook instance in Amazon Sagemaker.



2. Notebook Instance settings.

Notebook instance settings

Notebook instance name

Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

Notebook instance type

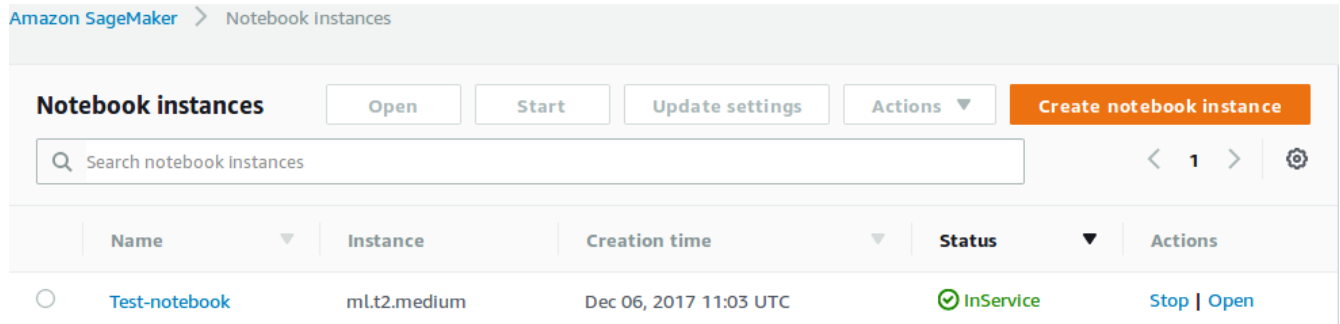
IAM role
Notebook instances require permissions to call other services including SageMaker and S3. Choose a role or let us create a role with the [AmazonSageMakerFullAccess](#) IAM policy attached.

VPC - optional
Notebook instances will have internet access independent of your VPC setting.

Subnet - optional

Security group(s) - optional

3. Notebook instance.

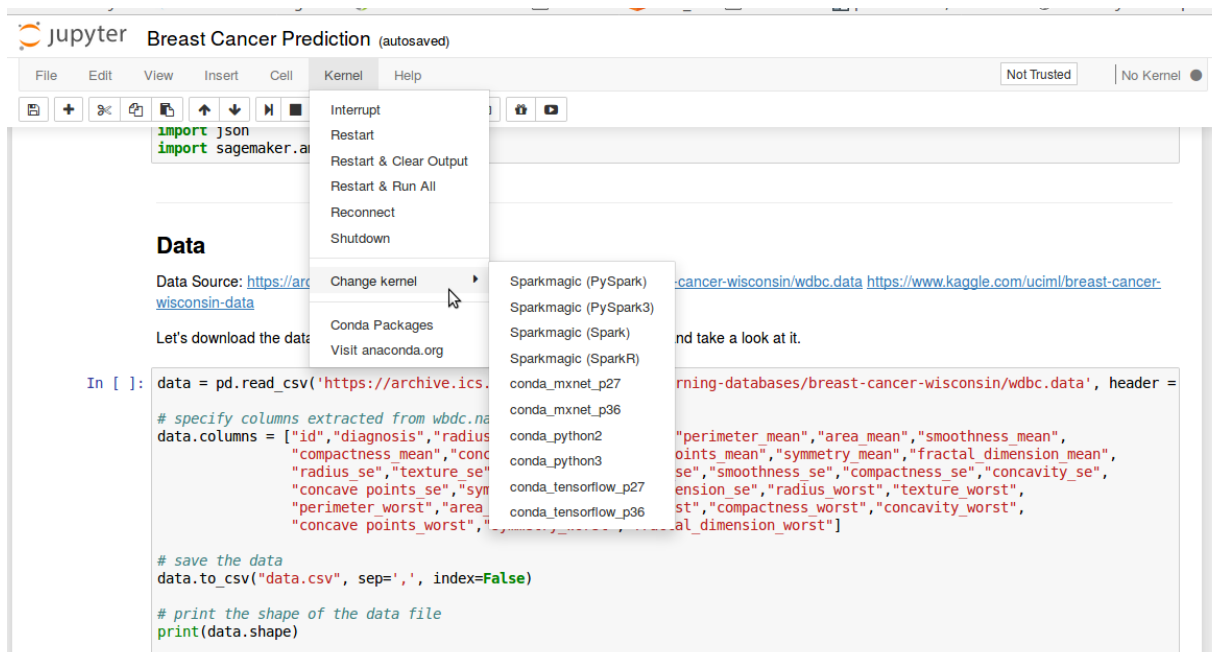


4. Open notebook instance.



5. Running sample notebook (breast cancer prediction).

5.1. Checking available kernels to run notebook.



6. Training ML model.

Train

Now we can begin to specify our linear model. Amazon SageMaker's Linear Learner actually fits many models in parallel, each with slightly different hyperparameters, and then returns the one with the best fit. This functionality is automatically enabled. We can influence this using parameters like:

- `num_models` to increase to total number of models run. The specified parameters will always be one of those models, but the algorithm also chooses models with nearby parameter values in order to find a solution nearby that may be more optimal. In this case, we're going to use the max of 32.
- `loss` which controls how we penalize mistakes in our model estimates. For this case, let's use absolute loss as we haven't spent much time cleaning the data, and absolute loss will be less sensitive to outliers.
- `wd` or `l1` which control regularization. Regularization can prevent model overfitting by preventing our estimates from becoming too finely tuned to the training data, which can actually hurt generalizability. In this case, we'll leave these parameters as their default "auto" though.

7. Specify containers images used for training.

Specify container images used for training and hosting SageMaker's linear-learner

```
In [15]: # See 'Algorithms Provided by Amazon SageMaker: Common Parameters' in the SageMaker documentation for an explanation of the
containers = {'us-west-2': '174872318107.dkr.ecr.us-west-2.amazonaws.com/linear-learner:latest',
             'us-east-1': '382416733822.dkr.ecr.us-east-1.amazonaws.com/linear-learner:latest',
             'us-east-2': '404615174143.dkr.ecr.us-east-2.amazonaws.com/linear-learner:latest',
             'eu-west-1': '438346466558.dkr.ecr.eu-west-1.amazonaws.com/linear-learner:latest'}
```

8. Creating hosting model for use.

```
In [18]: linear_hosting_container = {
          'Image': containers[boto3.Session().region_name],
          'ModelDataUrl': sm.describe_training_job(TrainingJobName=linear_job)['ModelArtifacts']['S3ModelArtifacts']
        }

create_model_response = sm.create_model(
    ModelName=linear_job,
    ExecutionRoleArn=role,
    PrimaryContainer=linear_hosting_container)

print(create_model_response['ModelArn'])

arn:aws:sagemaker:us-east-1:328115522647:model/linear-2017-12-06-11-27-45
```

9. Creating endpoint for hosting model with initial number of instances.

9.1. Configure endpoint.

```
In [19]: linear_endpoint_config = 'linear-endpoint-config-' + time.strftime("%Y-%m-%d-%H-%M-%S", time.gmtime())
print(linear_endpoint_config)
create_endpoint_config_response = sm.create_endpoint_config(
    EndpointConfigName=linear_endpoint_config,
    ProductionVariants=[{
        'InstanceType': 'ml.c4.2xlarge',
        'InitialInstanceCount': 1,
        'ModelName': linear_job,
        'VariantName': 'AllTraffic'}])

print("Endpoint Config Arn: " + create_endpoint_config_response['EndpointConfigArn'])

linear-endpoint-config-2017-12-06-11-35-56
Endpoint Config Arn: arn:aws:sagemaker:us-east-1:328115522647:endpoint-config/linear-endpoint-config-2017-12-06-11-35-56
```

9.2. Creating Endpoint.

```
In [20]: %%time

linear_endpoint = 'linear-endpoint-' + time.strftime("%Y%m%d%H%M", time.gmtime())
print(linear_endpoint)
create_endpoint_response = sm.create_endpoint(
    EndpointName=linear_endpoint,
    EndpointConfigName=linear_endpoint_config)
print(create_endpoint_response['EndpointArn'])

resp = sm.describe_endpoint(EndpointName=linear_endpoint)
status = resp['EndpointStatus']
print("Status: " + status)

sm.get_waiter('Endpoint_Created').wait(EndpointName=linear_endpoint)

resp = sm.describe_endpoint(EndpointName=linear_endpoint)
status = resp['EndpointStatus']
print("Arn: " + resp['EndpointArn'])
print("Status: " + status)

if status != 'InService':
    raise Exception('Endpoint creation did not succeed')

linear-endpoint-201712061140
arn:aws:sagemaker:us-east-1:328115522647:endpoint/linear-endpoint-201712061140
Status: Creating
Arn: arn:aws:sagemaker:us-east-1:328115522647:endpoint/linear-endpoint-201712061140
Status: InService
CPU times: user 132 ms, sys: 0 ns, total: 132 ms
Wall time: 14min 2s
```

10. Prediction invoking endpoint.

```
In [22]: runtime= boto3.client('sagemaker-runtime')

payload = np2csv(test_X)
response = runtime.invoke_endpoint(EndpointName=linear_endpoint,
                                   ContentType='text/csv',
                                   Body=payload)
result = json.loads(response['Body'].read().decode())
test_pred = np.array([r['score'] for r in result['predictions']])
```

```
In [24]: test_pred_class = (test_pred > 0.5)+0;
test_pred_baseline = np.repeat(np.median(train_y), len(test_y))

prediction_accuracy = np.mean((test_y == test_pred_class))*100
baseline_accuracy = np.mean((test_y == test_pred_baseline))*100

print("Prediction Accuracy:", round(prediction_accuracy,1), "%")
print("Baseline Accuracy:", round(baseline_accuracy,1), "%")

Prediction Accuracy: 97.8 %
Baseline Accuracy: 64.4 %
```

Results and training dataset are stored in S3 bucket. Custom built or other framework based algorithms can be hosted into container image and can be used like above. Sagemaker assumes S3 bucket already exists else you need to handle the exception of S3 bucket.

11. S3 Bucket output.

Viewing 1 to 3

<input type="checkbox"/>	Name	Last modified	Size	Storage class
<input type="checkbox"/>	linear-2017-12-06-11-27-45	--	--	--
<input type="checkbox"/>	train	--	--	--
<input type="checkbox"/>	validation	--	--	--

12. Hosted Model.

Models			
Create endpoint		Create endpoint configuration	Actions
<input type="text" value="Search models"/>		< 1 >	
<input type="radio"/>	Name	ARN	Creation time
<input type="radio"/>	linear-2017-12-06-11-27-45	arn:aws:sagemaker:us-east-1:328115522647:model/linear-2017-12-06-11-27-45	Dec 06, 2017 11:35 UTC

13. Endpoint.

Endpoint can be invoked from the console instead of jupyter notebook.

Endpoints

Update endpoint

Actions ▾

Create endpoint

🔍

Search endpoints

<

1

>

⚙️

	Name ▾	ARN	Creation time ▾	Status ▾	Last updated
<input type="radio"/>	linear-endpoint-201712061140	arn:aws:sagemaker:us-east-1:328115522647:endpoint/linear-endpoint-201712061140	Dec 06, 2017 11:40 UTC	✔️ InService	Dec 06, 2017 11:53 UTC

AWS Sagemaker endpoint can be embedded into your application by using or invoking the endpoint with help of boto3 and sagemaker SDK.

This document is simple take on AWS Sagemaker, it can be tuned accordingly to needs of the requirement at hand. It is essential to know docker to build docker files for your custom built or other frame work based algorithms.

References:

1. <https://github.com/awslabs/amazon-sagemaker-examples>
2. <https://docs.aws.amazon.com/sagemaker/latest/dg/whatis.html>
3. <https://pypi.python.org/pypi/sagemaker>